

Monte-Carlo examples in R

Serial vs parallel code

Meelis Utt

Setup

Let's source the setup (function, analytical solution, number of iterations).

```
source("Setup.R", echo = T)

##
## > options(scipen = 1000)
##
## > n <- 10000000
##
## > header <- c("n", "computational", "analytical", "error",
## +           "walltime", "type")
##
## > f <- function(x) {
## +   return(x^2 + x^4 + sin(x) + cos(x) + x^25)
## + }
##
## > analytical <- integrate(f, 0, 1)$value
```

Serial implementations

Let's start with a simple implementation of Monte-Carlo method.

```
MCser1 <- function(n){
  start <- Sys.time()
  i <- runif(n,0,1)
  EX <- mean(f(i))
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCser1"))
}
data.table(t(MCser1(n))) %>% setNames(header)

##           n      computational      analytical      error
## 1: 10000000 1.87312157698611 1.87296355073463 0.000158026251479537
##           walltime    type
## 1: 2.1086106300354 MCser1
```

Let's try a bit more vectorized solution, using the *apply function.

```
MCser2 <- function(n,ncols=1000){
  start <- Sys.time()
```

```

dt <- matrix(runif(n,0,1),ncol = ncols)
EX <- sapply(1:ncols,function(i,dt){
  EX <- dt[,i] %>% f %>% mean
},dt) %>% mean
end <- Sys.time()
time <- difftime(end,start)
error <- (EX - analytical) %>% abs
return(c(n,EX,analytical,error,time,"MCser2"))
}
data.table(t(MCser2(n))) %>% setNames(header)

##           n      computational      analytical      error
## 1: 10000000 1.87300128353283 1.87296355073463 0.0000377327982024056
##           walltime      type
## 1: 1.7653968334198 MCser2

```

Let's try an approach using data.table.

```

MCser3 <- function(n){
  start <- Sys.time()
  dt <- data.table(unif = runif(n,0,1))
  EX <- dt[,.(EX = mean(f(unif)))] %>% unlist %>% unname
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCser3"))
}
data.table(t(MCser3(n))) %>% setNames(header)

##           n      computational      analytical      error
## 1: 10000000 1.87298681647747 1.87296355073463 0.000023265742841172
##           walltime      type
## 1: 1.66933250427246 MCser3

```

Let's try divide-and-conquer approach with data.table.

```

MCser4 <- function(n,ncols=1000){
  start <- Sys.time()
  dt <- matrix(runif(n,0,1),ncol=ncols) %>% data.table
  EX <- dt[,lapply(.SD,f)][,lapply(.SD,mean)][,.(EX = sum(.SD)/ncols)] %>% unlist %>% unname
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCser4"))
}
data.table(t(MCser4(n))) %>% setNames(header)

##           n      computational      analytical      error
## 1: 10000000 1.87284836166252 1.87296355073463 0.000115189072112054
##           walltime      type
## 1: 1.81064772605896 MCser4

```

Parallel implementations

Let's try different parallel implementations. First let's start with package *parallel*.

```

MCpar1 <- function(n){
  start <- Sys.time()
  # Calculate the number of cores
  no_cores <- detectCores()
  # Initiate cluster
  cl <- makeCluster(no_cores)
  intermean <- parSapply(cl, rep(n/no_cores,no_cores),function(ni,f){
    EX <- mean(f(runif(ni,0,1)))
  },f)
  on.exit(stopCluster(cl))
  EX <- mean(intermean)
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCpar1"))
}
data.table(t(MCpar1(n))) %>% setNames(header)

```

```

##           n      computational      analytical      error
## 1: 10000000 1.87309074553834 1.87296355073463 0.000127194803714481
##           walltime      type
## 1: 1.22565364837646 MCpar1

```

Now let's try approach analogical to MCser2.

```

MCpar2 <- function(n){
  start <- Sys.time()
  # Calculate the number of cores
  no_cores <- detectCores()
  cl <- makeCluster(no_cores)
  dt <- matrix(runif(n,0,1),ncol = no_cores)
  intermean <- parSapply(cl, 1:no_cores,function(i,f,dt){
    EX <- mean(f(dt[,i]))
  },f,dt)
  on.exit(stopCluster(cl))
  EX <- mean(intermean)
  end <- Sys.time()
  error <- (EX - analytical) %>% abs
  time <- difftime(end,start)
  return(c(n,EX,analytical,error,time,"MCpar2"))
}
data.table(t(MCpar2(n))) %>% setNames(header)

```

```

##           n      computational      analytical      error
## 1: 10000000 1.87326255732499 1.87296355073463 0.000299006590361861
##           walltime      type
## 1: 3.25085091590881 MCpar2

```

This approach was not very good. But let's have one more try at analogical solution to MCser2.

```

MCpar2_2 <- function(n,ncols=1000){
  start <- Sys.time()
  # Calculate the number of cores
  no_cores <- detectCores()

```

```

cl <- makeCluster(no_cores)
dt <- matrix(runif(n,0,1),ncol = ncols)
intermean <- parSapply(cl, 1:ncols,function(i,f,dt){
  EX <- mean(f(dt[,i]))
},f,dt
)
on.exit(stopCluster(cl))
EX <- mean(intermean)
end <- Sys.time()
error <- (EX - analytical) %>% abs
time <- difftime(end,start)
return(c(n,EX,analytical,error,time,"MCpar2_2"))
}
data.table(t(MCpar2_2(n))) %>% setNames(header)

##           n      computational      analytical      error
## 1: 10000000 1.87314601528394 1.87296355073463 0.000182464549312922
##           walltime      type
## 1: 3.1351101398468 MCpar2_2

```

This solution was bit better, but still worse than the previous examples.

Let's try the package *foreach* now.

```

MCpar3 <- function(n){
  start <- Sys.time()
  # Calculate the number of cores
  no_cores <- detectCores()
  # Initiate cluster
  cl<-makeCluster(no_cores)
  # registerDoParallel(cl)
  EX <- foreach(ni = rep(n/no_cores,no_cores),.combine=mean,.export="f") %dopar%
    mean(f(runif(ni,0,1)))
  on.exit(stopCluster(cl))
  # stopImplicitCluster()
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCpar3"))
}
data.table(t(MCpar3(n))) %>% setNames(header)

## Warning: executing %dopar% sequentially: no parallel backend registered

##           n      computational      analytical      error
## 1: 10000000 1.87333968656737 1.87296355073463 0.000376135832744318
##           walltime      type
## 1: 1.82036471366882 MCpar3

```

Benchmarking

Now, let's visualize the walltimes of implemented solutions.

```
iterations <- c(2.5,5,7.5,10)*10**(5:7)
```

```
## Warning in c(2.5, 5, 7.5, 10) * 10^(5:7): longer object length is not a multiple
## of shorter object length

```

```

funs <- c(
  MCser1,MCser2,MCser3,MCser4,
  MCpar1,MCpar3 #MCpar2,MCpar2_2,
)
data <- sapply(funs,function(f,iterations){
  sapply(iterations,function(n){
    f(n)
  })
},iterations) %>%
matrix(ncol=6,byrow=T) %>%
data.table %>%
setNames(header) %>%
mutate_at(header[-grep(x=header,pattern="type")],as.numeric)

```

data

##		n	computational	analytical	error	walltime	type
##	1:	250000	1.873130	1.872964	0.00016643353	0.03385305	MCser1
##	2:	5000000	1.873030	1.872964	0.00006595650	0.69568539	MCser1
##	3:	75000000	1.872907	1.872964	0.00005654916	11.68683290	MCser1
##	4:	1000000	1.872530	1.872964	0.00043352101	0.14016461	MCser1
##	5:	250000	1.871317	1.872964	0.00164682859	0.20487976	MCser2
##	6:	5000000	1.872773	1.872964	0.00019064751	0.77347493	MCser2
##	7:	75000000	1.872913	1.872964	0.00005082639	11.00078845	MCser2
##	8:	1000000	1.873018	1.872964	0.00005454452	0.25182199	MCser2
##	9:	250000	1.875004	1.872964	0.00204025997	0.03408027	MCser3
##	10:	5000000	1.873260	1.872964	0.00029644140	0.65575528	MCser3
##	11:	75000000	1.873049	1.872964	0.00008524746	11.90426970	MCser3
##	12:	1000000	1.872933	1.872964	0.00003028198	0.13201976	MCser3
##	13:	250000	1.874320	1.872964	0.00135683934	0.23717046	MCser4
##	14:	5000000	1.873407	1.872964	0.00044341028	0.82303739	MCser4
##	15:	75000000	1.873138	1.872964	0.00017413568	11.54206014	MCser4
##	16:	1000000	1.873115	1.872964	0.00015100499	0.34342837	MCser4
##	17:	250000	1.871166	1.872964	0.00179786906	0.44379544	MCpar1
##	18:	5000000	1.873437	1.872964	0.00047358039	0.81169033	MCpar1
##	19:	75000000	1.872815	1.872964	0.00014894196	6.08725572	MCpar1
##	20:	1000000	1.873463	1.872964	0.00049947203	0.59535599	MCpar1
##	21:	250000	1.870885	1.872964	0.00207891484	0.46692729	MCpar3
##	22:	5000000	1.873080	1.872964	0.00011612905	1.09262633	MCpar3
##	23:	75000000	1.873205	1.872964	0.00024189381	10.45459032	MCpar3
##	24:	1000000	1.872843	1.872964	0.00012044145	0.57980013	MCpar3
##		n	computational	analytical	error	walltime	type

```

ggplot(data=data,aes(x=n,y=walltime,group=type,color=type)) +
  geom_point() +
  geom_line() +
  labs(
    title="Serial vs parallel implementations",
    x="Nr of iterations",
    y="Walltime (sec)"
  ) +
  theme(
    plot.title = element_text(hjust = 0.5),
    plot.subtitle = element_text(hjust = 0.5)
  )

```

```
) +  
guides(color=guide_legend(title="Function"))
```

