

Monte-Carlo examples in R

Serial vs parallel code

Meelis Utt

Setup

Let's source the setup (function, analytical solution, number of iterations).

```
source("Setup.R", echo = T)

##
## > options(scipen = 1000)
##
## > n <- 10000000
##
## > header <- c("n", "computational", "analytical", "error",
## +           "walltime", "type")
##
## > f <- function(x) {
## +   return(x^2 + x^4 + sin(x) + cos(x) + x^25)
## + }
##
## > analytical <- integrate(f, 0, 1)$value
```

Serial implementations

Let's start with a simple implementation of Monte-Carlo method.

```
MCser1 <- function(n){
  start <- Sys.time()
  i <- runif(n,0,1)
  EX <- mean(f(i))
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCser1"))
}
data.table(t(MCser1(n))) %>% setNames(header)

##           n      computational      analytical      error
## 1: 10000000 1.87353741604257 1.87296355073463 0.000573865307943233
##           walltime    type
## 1: 2.04202342033386 MCser1
```

Let's try a bit more vectorized solution, using the *apply function.

```
MCser2 <- function(n,ncols=1000){
  start <- Sys.time()
```

```

dt <- matrix(runif(n,0,1),ncol = ncols)
EX <- sapply(1:ncols,function(i,dt){
  EX <- dt[,i] %>% f %>% mean
},dt) %>% mean
end <- Sys.time()
time <- difftime(end,start)
error <- (EX - analytical) %>% abs
return(c(n,EX,analytical,error,time,"MCser2"))
}
data.table(t(MCser2(n))) %>% setNames(header)

```

```

##           n      computational      analytical      error
## 1: 10000000 1.87266314578146 1.87296355073463 0.000300404953171185
##           walltime    type
## 1: 1.60613822937012 MCser2

```

Let's try an approach using data.table.

```

MCser3 <- function(n){
  start <- Sys.time()
  dt <- data.table(unif = runif(n,0,1))
  EX <- dt[,.(EX = mean(f(unif)))] %>% unlist %>% unname
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCser3"))
}
data.table(t(MCser3(n))) %>% setNames(header)

```

```

##           n      computational      analytical      error
## 1: 10000000 1.87295819430205 1.87296355073463 0.00000535643257837393
##           walltime    type
## 1: 1.68406891822815 MCser3

```

Let's try divide-and-conquer approach with data.table.

```

MCser4 <- function(n,ncols=1000){
  start <- Sys.time()
  dt <- matrix(runif(n,0,1),ncol=ncols) %>% data.table
  EX <- dt[,lapply(.SD,f)][,lapply(.SD,mean)][,.(EX = sum(.SD)/ncols)] %>% unlist %>% unname
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCser4"))
}
data.table(t(MCser4(n))) %>% setNames(header)

```

```

##           n      computational      analytical      error
## 1: 10000000 1.87330731982586 1.87296355073463 0.00034376909123246
##           walltime    type
## 1: 1.88429856300354 MCser4

```

Parallel implementations

Let's try different parallel implementations. First let's start with package *parallel*.

```

MCpar1 <- function(n){
  start <- Sys.time()
  # Calculate the number of cores
  no_cores <- detectCores()
  # Initiate cluster
  cl <- makeCluster(no_cores)
  intermean <- parSapply(cl, rep(n/no_cores,no_cores),function(ni,f){
    EX <- mean(f(runif(ni,0,1)))
  },f)
  on.exit(stopCluster(cl))
  EX <- mean(intermean)
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCpar1"))
}
data.table(t(MCpar1(n))) %>% setNames(header)

```

```

##           n      computational      analytical      error
## 1: 10000000 1.87339102185438 1.87296355073463 0.000427471119751388
##           walltime    type
## 1: 1.32567048072815 MCpar1

```

Now let's try approach analogical to MCser2.

```

MCpar2 <- function(n){
  start <- Sys.time()
  # Calculate the number of cores
  no_cores <- detectCores()
  cl <- makeCluster(no_cores)
  dt <- matrix(runif(n,0,1),ncol = no_cores)
  intermean <- parSapply(cl, 1:no_cores,function(i,f,dt){
    EX <- mean(f(dt[,i]))
  },f,dt)
  on.exit(stopCluster(cl))
  EX <- mean(intermean)
  end <- Sys.time()
  error <- (EX - analytical) %>% abs
  time <- difftime(end,start)
  return(c(n,EX,analytical,error,time,"MCpar2"))
}
data.table(t(MCpar2(n))) %>% setNames(header)

```

```

##           n      computational      analytical      error
## 1: 10000000 1.87304020083651 1.87296355073463 0.000076650101882958
##           walltime    type
## 1: 4.93123769760132 MCpar2

```

This approach was not very good. But let's have one more try at analogical solution to MCser2.

```

MCpar2_2 <- function(n,ncols=1000){
  start <- Sys.time()
  # Calculate the number of cores
  no_cores <- detectCores()

```

```

cl <- makeCluster(no_cores)
dt <- matrix(runif(n,0,1),ncol = ncols)
intermean <- parSapply(cl, 1:ncols,function(i,f,dt){
  EX <- mean(f(dt[,i]))
},f,dt
)
on.exit(stopCluster(cl))
EX <- mean(intermean)
end <- Sys.time()
error <- (EX - analytical) %>% abs
time <- difftime(end,start)
return(c(n,EX,analytical,error,time,"MCpar2_2"))
}
data.table(t(MCpar2_2(n))) %>% setNames(header)

##           n      computational      analytical      error
## 1: 10000000 1.87279591153091 1.87296355073463 0.000167639203717762
##           walltime      type
## 1: 4.82105684280396 MCpar2_2

```

This solution was bit better, but still worse than the previous examples.

Let's try the package *foreach* now.

```

MCpar3 <- function(n){
  start <- Sys.time()
  # Calculate the number of cores
  no_cores <- detectCores()
  # Initiate cluster
  cl<-makeCluster(no_cores)
  # registerDoParallel(cl)
  EX <- foreach(ni = rep(n/no_cores,no_cores),.combine=mean,.export="f") %dopar%
    mean(f(runif(ni,0,1)))
  on.exit(stopCluster(cl))
  # stopImplicitCluster()
  end <- Sys.time()
  time <- difftime(end,start)
  error <- (EX - analytical) %>% abs
  return(c(n,EX,analytical,error,time,"MCpar3"))
}
data.table(t(MCpar3(n))) %>% setNames(header)

## Warning: executing %dopar% sequentially: no parallel backend registered

##           n      computational      analytical      error
## 1: 10000000 1.87332913732254 1.87296355073463 0.000365586587907529
##           walltime      type
## 1: 2.03123497962952 MCpar3

```

Benchmarking

Now, let's visualize the walltimes of implemented solutions.

```

iterations <- c(1,2.5,5,7.5)*10**(7)#(5:7)
funs <- c(
  MCser1,MCser2,MCser3,MCser4,

```

```

MCpar1,MCpar3 #MCpar2,MCpar2_2,
)
data <- sapply(funs,function(f,iterations){
  sapply(iterations,function(n){
    f(n)
  })
},iterations) %>%
matrix(ncol=6,byrow=T) %>%
data.table %>%
setNames(header) %>%
mutate_at(header[-grep(x=header,pattern="type")],as.numeric)

```

data

##		n	computational	analytical	error	walltime	type
##	1:	10000000	1.873144	1.872964	0.00018006442	1.679950	MCser1
##	2:	25000000	1.872934	1.872964	0.00002993648	4.520063	MCser1
##	3:	50000000	1.872776	1.872964	0.00018749121	9.524911	MCser1
##	4:	75000000	1.873037	1.872964	0.00007335057	14.163665	MCser1
##	5:	10000000	1.873000	1.872964	0.00003679440	1.872934	MCser2
##	6:	25000000	1.873036	1.872964	0.00007210419	4.387838	MCser2
##	7:	50000000	1.872938	1.872964	0.00002604306	8.864731	MCser2
##	8:	75000000	1.872995	1.872964	0.00003098338	12.007296	MCser2
##	9:	10000000	1.872698	1.872964	0.00026519163	1.755220	MCser3
##	10:	25000000	1.872981	1.872964	0.00001748658	4.360301	MCser3
##	11:	50000000	1.872957	1.872964	0.00000608254	9.454168	MCser3
##	12:	75000000	1.872898	1.872964	0.00006552740	13.830856	MCser3
##	13:	10000000	1.873285	1.872964	0.00032096333	2.406975	MCser4
##	14:	25000000	1.872999	1.872964	0.00003529201	4.818219	MCser4
##	15:	50000000	1.872987	1.872964	0.00002313410	9.413631	MCser4
##	16:	75000000	1.872992	1.872964	0.00002803494	12.982518	MCser4
##	17:	10000000	1.872892	1.872964	0.00007109093	1.288364	MCpar1
##	18:	25000000	1.873019	1.872964	0.00005506899	2.609506	MCpar1
##	19:	50000000	1.872944	1.872964	0.00001965172	4.578090	MCpar1
##	20:	75000000	1.872952	1.872964	0.00001173007	6.752777	MCpar1
##	21:	10000000	1.872822	1.872964	0.00014118409	2.141526	MCpar3
##	22:	25000000	1.873535	1.872964	0.00057168676	4.074641	MCpar3
##	23:	50000000	1.872938	1.872964	0.00002530814	7.826763	MCpar3
##	24:	75000000	1.872941	1.872964	0.00002210452	11.566591	MCpar3
##		n	computational	analytical	error	walltime	type

```

ggplot(data=data,aes(x=n,y=walltime,group=type,color=type)) +
  geom_point() +
  geom_line() +
  labs(
    title="Serial vs parallel implementations",
    x="Nr of iterations",
    y="Walltime (sec)"
  ) +
  theme(
    plot.title = element_text(hjust = 0.5),
    plot.subtitle = element_text(hjust = 0.5)
  ) +
  guides(color=guide_legend(title="Function"))

```

Serial vs parallel implementations

